

AI_AGENT

Multi-Agent AI Governance Evaluation System

Evaluating AI agents across Standards · Safety · Assurance


Tongye Ding, Andrew Shi, Siwen Shen

Capstone Project | Sponsor: UNICC / Dr. Andrés Fortino (NYU)

GOVERN-
ANCE
EVAL.
AGENT

SYSTEM ARCHITECTURE

What We Built



WHAT WE BUILT

AI_AGENT is an end-to-end multi-agent system that evaluates AI use cases for governance readiness before deployment in the UNICC AI Sandbox. Users submit a use case description and the system automatically runs it through three specialized AI agents in parallel, then synthesizes a final binding decision.



Standards Agent

Checks compliance against EU AI Act (2024), US AI Bill of Rights (2023), IEEE 7001/7003, ISO/IEC 23894, and UNESCO AI Ethics — identifying gaps and required controls per standard.



Safety Dimensions Agent

Scores six risk dimensions: Harmfulness, Bias & Fairness, Transparency, Deception, Privacy, and Accountability — each from 1–5, visualized in a bar chart on the dashboard.



Assurance & Accountability Agent

Evaluates auditability, human oversight mechanisms, monitoring/logging, and risk control evidence — producing structured gaps and recommended assurance controls.

Stack: React + FastAPI + Llama 3.2 (Ollama) · Local SLM — no external APIs · Immutable audit log (JSON)

THE TOOL IN ACTION

INPUT · Submit an AI Use Case

2026/4/11 16:10 frontend

AI_AGENT Governance Dashboard

Multi-agent evaluation across standards, safety dimensions, and assurance controls.

Standards
EU AI Act, ISO, IEEE, UNESCO

Safety
Bias, privacy, transparency, accountability

Assurance
Audit trails, controls, fallback readiness

Current system decision

Final output

Flag for Review

Submit a use case to get the combined governance judgment.

Evaluate an AI use case

Use case title
AI Hiring Screening System

Description
An AI system ranks job applicants using resume text, application history, and screening question responses before human review.

Intended users
HR recruiters and talent ac

Affected groups
Job applicants, including p

Data sources
Resume text, application h

Deployment context
Used internally by a hiring team to prioritize candidates before final human review.

Current controls

Cross-agent risks

Harmfulness
Bias & Fairness

Privacy

Cross-agent gaps

Lack of transparency in decision-making process

Insufficient consideration for diverse applicant groups

Insufficient data protection measures

Lack of transparency in AI decision-making processes

Insufficient testing and validation of the system's fairness and

Cross-agent controls

Implement model interpretability techniques

Provide clear explanations for bias and assumptions

Conduct diversity and inclusion audits

Establish clear guidelines for handling sensitive information

Implement robust data encryption and access controls

Implement diversity and inclusion

localhost:5173 1/4

- Use case title, description, intended users
- Affected groups, data sources, deployment context
- Sensitive domain flag + Human oversight toggle

OUTPUT · Governance Summary Results

2026/4/11 16:25 frontend

Standards Compliance Agent

Executive summary
The proposed Customer Support FAQ Chatbot aligns with several standards but has gaps in data governance and human oversight.

Triggered standards
EU AI Act (2024)
US AI Bill of Rights (2023)
IEEE 7001 ISO/IEC 23894
UNESCO Recommendation on the Ethics of AI

Key gaps

- Data governance: Lack of clear data sources and provenance
- Human oversight: Insufficient controls for complex queries

Required controls

- Clear data governance policies
- Enhanced human oversight mechanisms for complex queries

Recommendation
Approve with Conditions

Safety Dimensions Agent

Executive summary
Customer Support FAQ Chatbot uses AI to answer frequently asked questions, with some potential risks identified.

Recommendation
Approve with Conditions

Assurance & Accountability Agent

Executive summary
The proposed Customer Support FAQ Chatbot has a solid foundation for deployment, but some gaps in evidence and accountability structure need attention.

Evidence gaps

- Lack of testing data on edge cases
- Insufficient documentation on fallback procedures

Assurance controls

- Implement automated testing to validate edge cases and ensure the chatbot's accuracy.
- Develop a comprehensive fallback procedure for handling complex queries, including escalation protocols for human support teams.
- Regularly review and update the knowledge base content to maintain its relevance and accuracy.

Recommendation
Approve with Conditions

localhost:5173 3/3

- Cross-agent risks, gaps & controls synthesized
- Per-agent executive summaries + recommendations
- Safety dimension bar chart · Final governance verdict

Questions & Discussion

- ⚖️ Standards: EU AI Act · US AI Bill of Rights · IEEE · ISO · UNESCO
- 🛡️ Safety: 6 risk dimensions scored — Harmfulness, Bias, Privacy, Transparency...
- ✓ Assurance: Audit trails, human oversight, monitoring controls
- ⚡ Stack: React · FastAPI · Llama 3.2 via Ollama · 100% on-premises